



# PWNIES INTRO WORKSHOP

## Day 1: Shellcode

Department of Computer Science, University of Copenhagen

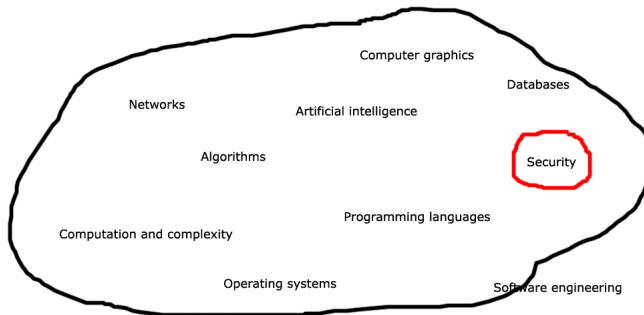
November 11, 2013

Slide 1/42




# Computer and information security

## COMPUTER SCIENCE



# ZOOM; ENHANCE

## COMPUTER SECURITY




Many of these are common categories in CTFs.



## ZOOM; ENHANCE

## COMPUTER SECURITY



Our focus in this workshop.



# This week's program

- Shellcode and x86 assembly
- Reverse engineering and debugging
- Buffer overflow exploits
- Network-based exploits
- Advanced topics (format string, return-oriented programming)



# Today's program

- 0 Hardware and lowlevel programming
- 1 CPU, x86 assembler instructions, and shellcode
- 2 Memory and the stack
- 3 Improving your shellcode
- 4 Extra exercises




# Today's program

## 0 Hardware and lowlevel programming

- 1 CPU, x86 assembler instructions, and shellcode
- 2 Memory and the stack
- 3 Improving your shellcode
- 4 Extra exercises



## 0. Hardware and lowlevel programming



At the base level, computers are just electronic components


...







# 0. Hardware and lowlevel programming



...and mapped to 0's and 1's.



# 0. Hardware and lowlevel programming

- Everything the computer does is *really just numbers*.



## 0. Hardware and lowlevel programming

- Everything the computer does is *really just numbers*.
- Example: In the MIPS architecture, the sequence 0000000010101111000000000100000 is an instruction for the computer to add two numbers.



## 0. Hardware and lowlevel programming

- Everything the computer does is *really just numbers*.
- Example: In the MIPS architecture, the sequence  
00000000101011111000000000100000 is an instruction  
for the computer to add two numbers.
- What if we change *one* bit,  
01000000101011111000000000100000?



## 0. Hardware and lowlevel programming

- Everything the computer does is *really just numbers*.
- Example: In the MIPS architecture, the sequence 0000000010101111000000000100000 is an instruction for the computer to add two numbers.
- What if we change *one* bit,  
0100000010101111000000000100000?
- The meaning changes entirely - now it's a **jump** instruction.



## 0. Hardware and lowlevel programming

- Everything the computer does is *really just numbers*.
- Example: In the MIPS architecture, the sequence 0000000010101111000000000100000 is an instruction for the computer to add two numbers.
- What if we change *one* bit, 0100000010101111000000000100000?
- The meaning changes entirely - now it's a **jump** instruction.
- Going deeper: Binary sequences also represent numbers (the jump one is 1085243424), ascii values (@&#175;&#8364;), and more.



# 0. Hardware and lowlevel programming

- A CPU works by executing a flow of instructions.





# 0. Hardware and lowlevel programming

- A CPU works by executing a flow of instructions.
- Sometimes, a CPU can be 'tricked' into executing your instructions instead of the original program.



# 0. Hardware and lowlevel programming

- A CPU works by executing a flow of instructions.
- Sometimes, a CPU can be 'tricked' into executing your instructions instead of the original program.
- Today, you'll learn how to write these instructions in the Intel x86 architecture.



# Today's program

- 0 Hardware and lowlevel programming
- 1 **CPU, x86 assembler instructions, and shellcode**
- 2 Memory and the stack
- 3 Improving your shellcode
- 4 Extra exercises



# 1.0 CPU, x86 assembler instructions, and shellcode

- A CPU can only execute simple instructions like `add` or `mov`.



# 1.0 CPU, x86 assembler instructions, and shellcode

- A CPU can only execute simple instructions like `add` or `mov`.
- Usually one line of c code will translate into several lines of assembler.



## 1.1 registers

- A register is a 32 bit memory inside the CPU.



## 1.1 registers

- A register is a 32 bit memory inside the CPU.
- EAX, EBX, ECX, EDX (EDI, ESI) are general purpose registers.



## 1.1 registers

- A register is a 32 bit memory inside the CPU.
- EAX, EBX, ECX, EDX (EDI, ESI) are general purpose registers.
- ESP, EBP are used for memory management.





## 1.1 registers

- A register is a 32 bit memory inside the CPU.
- EAX, EBX, ECX, EDX (EDI, ESI) are general purpose registers.
- ESP, EBP are used for memory management.
- EIP points to the next instruction.



## 1.1 registers

- A register is a 32 bit memory inside the CPU.
- EAX, EBX, ECX, EDX (EDI, ESI) are general purpose registers.
- ESP, EBP are used for memory management.
- EIP points to the next instruction.
- EFLAGS will not be discussed.



# 1.1 registers

32	24	16	8
E_X			
		_X	
		_H	_L



## 1.2 Example: mov instruction:

mov a, b

a = b

- mov register, value  
moves value (like 5) into register (like eax)
- mov register a, register b  
moves the content of reg. b into reg. a
- mov register a, [pointer]  
moves the content on the address "pointer" into reg. a



## 1.3 syscall:

Stops the instruction execution, and lets the operating system execute a function.

<http://syscalls.kernelgrok.com/>

- eax contains the syscall number.
- ebx, ecx and edx contains the arguments.
- 0x80 makes the call.



## 1.4 compiling, linking, and executing a program:

```
compile:  nasm -f elf32 <name>.asm  
link:     ld -melf_i386 -o <name> <name>.o  
run       ./<name>
```

Replace **<name>** with the filename of your program.



## 1.5 x86 assembler template

```
[bits 32]  
section .text  
global _start
```

```
_start:  
mov eax, str
```

```
section .data  
str: db 'lolhej', 0x0
```



-----





# Today's program

- 0 Hardware and lowlevel programming
- 1 CPU, x86 assembler instructions, and shellcode
- 2 **Memory and the stack**
- 3 Improving your shellcode
- 4 Extra exercises



## 2. Memory and the stack

Memory on x86 is very complex. Long story short:

- Memory ranges from `0x00000000` to `0xffffffff`.



## 2. Memory and the stack

Memory on x86 is very complex. Long story short:

- Memory ranges from `0x00000000` to `0xffffffff`.
- Every address contains 1 byte, but are normally used in chunks of 4 bytes (a dword).



## 2. Memory and the stack

Memory on x86 is very complex. Long story short:

- Memory ranges from `0x00000000` to `0xffffffff`.
- Every address contains 1 byte, but are normally used in chunks of 4 bytes (a dword).
- x86 is little endian (the 'least significant byte' is saved first). The number `0xDECAFEBAD` consists of `0xAD`, `0xFB`, `0xCA`, `0xDE`, *in that order*!



## 2. Memory and the stack

Memory on x86 is very complex. Long story short:

- Memory ranges from `0x00000000` to `0xffffffff`.
- Every address contains 1 byte, but are normally used in chunks of 4 bytes (a dword).
- x86 is little endian (the 'least significant byte' is saved first). The number `0xDECAFEBAD` consists of `0xAD`, `0xFB`, `0xCA`, `0xDE`, *in that order*!
- The addresses are local for your program.



## 2. Memory and the stack

Example:

0x000c	...
0x0008	...
0x0004	0xDECAFBAD
0x0000	...



## 2. Memory and the stack

Example:

0x000c	...
0x0008	...
0x0004	0xDECAFBAD
0x0000	...


Zooming in at address 0x0004:

0x0007	0xDE
0x0006	0xCA
0x0005	0xFB
0x0004	0xAD



## 2. Memory and the stack

x86 memory is divided into segments:





## 2. Memory and the stack

- We have made some assembler code which spawns a shell.



## 2. Memory and the stack

- We have made some assembler code which spawns a shell.
- Problem: We are using labels (why is this a problem?).



## 2. Memory and the stack

- We have made some assembler code which spawns a shell.
- Problem: We are using labels (why is this a problem?).
- Solution: We have plenty of space in dynamic memory! It's time to look at the stack.



## 2. Memory and the stack

- The stack is a 'first in, last out' structure that grows *downwards*.



## 2. Memory and the stack

- The stack is a 'first in, last out' structure that grows *downwards*.
- ESP, the stack pointer, points to the top of the stack.



## 2. Memory and the stack

- The stack is a 'first in, last out' structure that grows *downwards*.
- ESP, the stack pointer, points to the top of the stack.
- The **push** instruction adds an element (4 bytes) to the top of the stack and decrements ESP.




## 2. Memory and the stack

- The stack is a 'first in, last out' structure that grows *downwards*.
- ESP, the stack pointer, points to the top of the stack.
- The **push** instruction adds an element (4 bytes) to the top of the stack and decrements ESP.
- The **pop** instruction increments ESP by 4 (effectively removing the 4 bytes at the top).



## 2. Memory and the stack

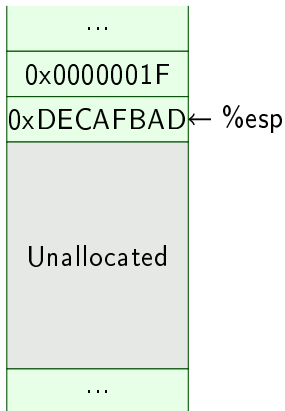


```
-> push 0xDECAFBAD  
    push 'hey'  
    push byte 0x42  
    pop  eax
```





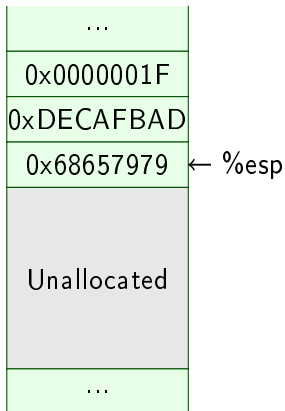
## 2. Memory and the stack



```
-> push 0xDECAFBAD
    push 'hey'
    push byte 0x42
    pop eax
```



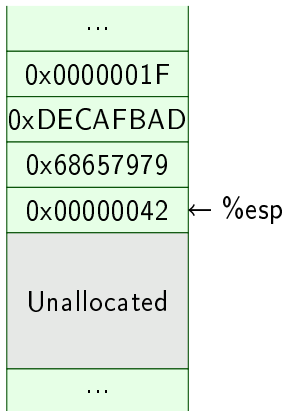
## 2. Memory and the stack



```
push 0xDECAFBAD
push 'heyy' ('h' = 0x68. Check man ascii.)
-> push byte 0x42
pop eax
```




## 2. Memory and the stack



```
push 0xDECAFBAD
push 'hey'
push byte 0x42
-> pop eax
```



## 2. Memory and the stack



```
push 0xDECAFBAD
push 'heyy'
push byte 0x42
pop  eax (EAX now contains 0x42)
```



## 2. Memory and the stack

```
-----
| Exercise: Write a program that executes |
|           /bin/sh without using labels! |
| Hint:      Strings are null-terminated. |
|-----
```

```

\
(\^ )
o 0\_____ -
\_ /          \ \
  \ _____ / \ \
  / \  | | \ \  \ \
  | | \ | |  \ \

```



# Today's program

- 0 Hardware and lowlevel programming
- 1 CPU, x86 assembler instructions, and shellcode
- 2 Memory and the stack
- 3 Improving your shellcode**
- 4 Extra exercises



### 3. Improving your shellcode

- `xor eax, eax`
  - set `eax` to 0.



### 3. Improving your shellcode

- `xor eax, eax`
  - set `eax` to 0.
- `mul ecx`
  - `edx:eax = eax*ecx`





### 3. Improving your shellcode

- `xor eax, eax`
  - set `eax` to 0.
- `mul ecx`
  - `edx:eax = eax*ecx`
- `mov al, 5`
  - move 5 to the least significant byte in `eax`.



### 3. Improving your shellcode

```

-----
| Exercises:                                     |
| a. Write shellcode that executes /bin/sh |
|      without null bytes!                     |
| b. Optimize your shellcode to using as      |
|      few bytes as possible!                   |
|-----

```

```

\
(\^)
o 0\-----
\_/      \
 \      /\
 /\  ||\ \
 ||\ || \

```



# Today's program

- 0 Hardware and lowlevel programming
- 1 CPU, x86 assembler instructions, and shellcode
- 2 Memory and the stack
- 3 Improving your shellcode
- 4 **Extra exercises**



## 4. Extra exercises!

```

-----
| You've found a vulnerability in a service |
| run by <insert evil>, allowing you to    |
| execute shellcode.                       |
| You're interested in orienting yourself  |
| to find potentially interesting files...  |
|                                           |
| Write shellcode that executes /bin/ls with |
| '-l' as argument!                       |
|                                           |
-----

```

```

\
(\^ )
o 0\_____ -
\_/          \
 \      /\
//\  ||\ \ \
||\ \ || \ \

```



## 4. Extra exercises!

```

-----
| Huh. It seems they have a file called      |
| gold.txt lying around. Let's check it out. |
|                                             |
| Write shellcode that writes the contents   |
| of a text file to standard output!        |
| Hint: sys_open, sys_read, sys_write.      |
|           For testing, create your own gold.txt |
|           file.                            |
-----

```

```

\
(\^)
o 0\_____ -
\_/          \
 \    ____  /\
  /\  ||\ \  \
    ||\ \ || \
    ||\ \ || \

```



## 4. Extra! Compare instruction

`cmp a, b`

`a ? b`

compares `a` and `b`, the result is stored in a special register, and can be used by instructions like `jeq`, `jne`, `jge`, `jle`

- `cmp register, value`  
compares the value of register and value
- `cmp register a, register b`  
compares the value register `a` and `b`
- `add register a, [register b + x]`  
compares the value of register `a` and the value on the address of register `b + x`



## 4. Extra! Conditional jump instruction

j\_\_ label

Used after compare, if the condition is true then the instruction pointer is moved to the address given by label.

- ja: jump if above
- jb: jump if below
- jeq: jump if equal
- jne: jump if not equal
- jae: jump if above or equal
- jbe: jump if below or equal
- jmp: always jump (does not have to be after a compare)

Compare and conditional jump instructions can be used to form loops.



## 4. Extra exercises!

- ```

-----
| a. Write a program that calculates the      |
| factorial of some number n.                |
| (For now, hardcode the number n in your    |
| program. We'll get to function calls      |
| another day.)                             |
| b. Write a program that calculates the nth  |
| fibonacci number for some number n.       |
|
-----

```

```

\
(\^ )
o 0\-----
\_/      \
 \      /\
  /\  ||\  \
  ||\||  \

```

